

# Machines That Understand: A Theoretical Survey of Deep Learning's Role in Shaping Natural Language Processing

Vasisht Boddula

Independent Researcher

Alumnus, Department of Computer Science, Aurora's Technological and Research Institute, India

Published: April 16, 2025

---

## Abstract

Natural Language Processing (NLP) has evolved from rule-based systems - relying on manually encoded grammar rules - to statistical models using probabilistic techniques such as **n-gram modelling** and **Hidden Markov Models (HMMs)**. These approaches introduced data-driven learning but lacked semantic richness and contextual depth. As tasks grew more complex, traditional systems struggled with ambiguity and scalability across domains. These limitations paved the way for deep learning, which offered solutions through automatic representation learning. This paper retraces this evolution to show how contemporary NLP systems now leverage contextualized models and neural computation to go beyond surface-level language processing and closer toward language understanding.

The introduction of deep learning brought advanced architectures capable of modelling language in nonlinear and contextual ways. Methods like word embeddings, **Recurrent Neural Networks (RNNs)**, and later, Transformer architectures allowed models to capture long-term dependencies and bidirectional context. The Transformer, using a self-attention mechanism, enabled large-scale pretraining, leading to the development of models like **Bidirectional Encoder Representations from Transformers (BERT)**, **Generative Pre-trained Transformer (GPT)**, and **Text-To-Text Transfer Transformer (T5)**. These task-agnostic models dramatically reduced the need for handcrafted features and ushered in a shift from traditional pipeline systems to unified frameworks capable of language understanding and generation.

Modern NLP depends heavily on semantic modelling—learning the meaning of words through their context in large corpora. Unlike earlier approaches, which were sensitive to ambiguity and domain shifts, models like GPT and BERT dynamically learn meaning across tasks and languages. This flexibility supports applications such as machine translation, summarization, and question answering. These systems are now actively used in domains including healthcare (clinical document summarization), legal technology (document review), finance (sentiment-based forecasting), and education (intelligent tutoring systems). This paper explores how these semantic advancements enable machines to adapt to real-world problems effectively.

While deep learning has propelled NLP forward, it introduces concerns regarding interpretability, fairness, and ethical deployment. Large-scale models function as “black boxes,” making their decision logic hard to explain. Moreover, these models often reflect societal biases found in training data - impacting fairness in critical domains. Additional risks include misinformation (incorrect yet plausible outputs), hallucination (generation of non-factual content), and data leakage (memorization of sensitive inputs). This paper discusses the need for responsible AI practices such as fairness-aware

training, transparency mechanisms like model cards, and regulation to build robust, explainable, and socially aligned NLP systems for the future.

Natural Language Processing continues to advance toward adaptability, multimodality, and low-resource generalization. Hybrid models combining deep learning and symbolic logic offer greater explainability, while few-shot prompting reduces data dependence. Research now emphasizes multilingual fairness, domain-specific transferability, and sustainable training practices to mitigate ecological and economic costs. With ongoing developments in multimodal Artificial Intelligence - integrating language with vision, sound, and structured knowledge - machines are beginning to understand language more like humans do. This paper concludes by outlining a theoretical roadmap for NLP research focused on semantic precision, inclusive design, and cognitive alignment, paving the way for ethically grounded, intelligent language systems.

**Keywords:** Natural Language Processing (NLP), Deep Learning, Transformer Models, BERT, GPT, Semantic Modelling, Neural Architectures, Contextual Embeddings, Language Understanding, Ethical AI

## 1. Introduction

Natural Language Processing (NLP) is a key area of artificial intelligence that enables machines to understand and generate human language. It has evolved from rule-based systems with handcrafted grammar to statistical models using probabilistic reasoning. While early approaches helped automate language tasks, they often failed to address ambiguity, scale, and cross-domain variation effectively.

The emergence of statistical NLP introduced models like n-grams and Hidden Markov Models (HMMs), which allowed systems to learn from text data. However, these models had limited capacity for capturing long-term dependencies and semantic meaning. This led to the adoption of deep learning methods, which offered better generalization through layered representations and context modelling.

**Recurrent Neural Networks (RNNs), Long Short-Term Memory (LSTM) networks, and Transformer architectures** marked a significant advancement. Models such as **Bidirectional Encoder Representations from Transformers (BERT), Generative Pre-trained Transformer (GPT), and Text-To-Text Transfer Transformer (T5)** demonstrated the power of pretraining on large corpora, enabling models to perform multiple language tasks with minimal fine-tuning.

Despite these advancements, modern NLP systems raise concerns related to interpretability, bias, misinformation, and ethical use. As these models are increasingly used in critical areas such as healthcare, education, and law, it becomes essential to ensure transparency and accountability.

This paper offers a theoretical overview of NLP's evolution, focusing on deep learning paradigms and their impact on language understanding. It explores the transition from foundational methods to general-purpose models, examines key architectural trends, and discusses limitations and future research directions in building more responsible and scalable language systems.

## 2. Background and Evolution of Natural Language Processing

The growth of Natural Language Processing (NLP) reflects decades of research shaped by new ideas about how machines can process human language. In its early years, NLP was built on **rule-based systems**, where experts wrote grammar rules by hand to help machines understand language. These

systems were precise but rigid. They worked well in limited scenarios, but they could not easily adapt to new languages, informal expressions, or complex sentence structures.

Later, as more digital text became available, researchers began using **statistical methods** to model language. This shift led to **statistical NLP**, where models learned patterns from large text datasets instead of relying on fixed rules. Techniques like **n-gram models** and **Hidden Markov Models (HMMs)** allowed systems to predict word sequences and classify text based on probability. These models improved flexibility, but they still struggled with deep meaning and long-term context. For example, they could identify frequent patterns but could not explain what a sentence *really meant*.

The next breakthrough came with **deep learning**, which allowed machines to learn complex patterns and hidden meanings in language. Early neural models like **Recurrent Neural Networks (RNNs)** and **Long Short-Term Memory (LSTM)** networks helped machines understand sequences of words over time. However, they were still limited in how much context they could remember. That changed with the **Transformer architecture**, which introduced **attention mechanisms**. Transformers allow models to consider all words in a sentence at once, capturing deeper relationships and meaning more effectively.

Using Transformers, researchers created **pretrained language models** such as **BERT (Bidirectional Encoder Representations from Transformers)**, **GPT (Generative Pre-trained Transformer)**, and **T5 (Text-To-Text Transfer Transformer)**. These models are trained on large amounts of text and can then be adapted to many different tasks with minimal changes. Instead of building a new model for every task, developers can now fine-tune a single powerful model for tasks like summarization, translation, or question answering.

The history of NLP shows a clear shift—from rule-based systems, to statistical learning, to deep neural networks. Each stage has brought improvements in accuracy and flexibility, but also new challenges. Today's models are more powerful than ever, but they are also harder to understand and control. As we move forward, it is important to study both the strengths and limitations of these systems, and to ask deeper questions about what it means for machines to truly understand human language.

### 3. Deep Learning Architectures in Natural Language Processing

The use of **deep learning** has completely changed how machines understand language. Instead of relying on fixed rules or simple patterns, deep learning allows models to learn directly from large amounts of text. These models can discover hidden relationships between words, understand context, and make more accurate predictions. This section explains the key architectures that have made deep learning successful in Natural Language Processing (NLP).

The first major architecture used in NLP was the **Recurrent Neural Network (RNN)**. RNNs were designed to work with sequences of data, such as sentences. They process words one at a time and use memory to connect the current word with previous words. However, RNNs had trouble remembering information from long sentences, and training them was often slow and unstable. To solve this, researchers developed **Long Short-Term Memory (LSTM)** networks. LSTMs improved memory and made it easier for models to handle longer sentences, but they still had limits when it came to learning from large datasets.

A breakthrough came with the **Transformer** architecture. Unlike RNNs, Transformers do not process words in order. Instead, they use an **attention mechanism** to look at all the words in a sentence at the same time. This allows the model to understand how words relate to each other, even if they are far

apart in the sentence. The attention mechanism is what gives Transformers their power—it helps models focus on the most important parts of the input.

Transformers also made it possible to train very large models on huge datasets. This led to the development of **pretrained language models**, which are trained once on general language data and then fine-tuned for specific tasks. Some of the most well-known models include:

- **BERT (Bidirectional Encoder Representations from Transformers)**: learns meaning from both left and right context in a sentence.
- **GPT (Generative Pre-trained Transformer)**: is trained to predict the next word and is strong at generating text.
- **T5 (Text-To-Text Transfer Transformer)**: turns all tasks—such as translation or classification—into a text generation problem.

These models are not designed for a single task. Instead, they can be used for many different tasks, such as summarization, sentiment analysis, question answering, or translation, often with little extra training. This is a big change from earlier NLP systems, which required separate models for each task.

In summary, deep learning architectures have made NLP more flexible, more accurate, and easier to scale. They allow machines to understand not just individual words, but also the context and meaning behind them. However, as these models become more powerful, they also become more complex and harder to explain. The next sections will explore how these models are used in real-world applications and the challenges they bring.

#### 4. Semantic Modelling and Language Understanding

Understanding language is more than recognizing words in the correct order - it involves capturing meaning, context, and relationships between ideas. This process is known as **semantic modelling**. In Natural Language Processing (NLP), semantic modelling refers to how machines learn the meaning of words, phrases, and sentences, and how they use this meaning to perform tasks like translation, summarization, or answering questions.

Traditional models in NLP treated words as isolated units. They used one-hot encoding, where each word had a unique position in a list, but this did not show how words are related. For example, the words “doctor” and “nurse” would look completely different to the machine, even though they are related in meaning. Deep learning changed this by introducing **word embeddings**, which represent words as vectors in a space where similar words are closer together. This allowed machines to understand relationships like “king is to queen as man is to woman.”

More advanced models like **BERT (Bidirectional Encoder Representations from Transformers)** and **GPT (Generative Pre-trained Transformer)** go further. They do not just look at a single word - they look at all the words around it to understand its meaning in context. For example, the word “bank” can mean a riverbank or a financial institution, depending on the sentence. These models use **contextual embeddings** to decide which meaning is correct based on surrounding words. This is a major improvement in how machines understand human language.

Semantic modelling also helps in more complex tasks. In **question answering**, the model must understand what the question is asking and find the correct answer in a text. In **summarization**, it must identify the most important ideas. In **machine translation**, it must understand the meaning of the sentence before converting it into another language. In all these cases, understanding context and meaning is essential, and that is what modern NLP models are designed to do.

These improvements have allowed NLP systems to be used in real-world fields like **healthcare, law, and education**. For instance, a model can now summarize a medical report, extract legal terms from a contract, or assist in grading essays by understanding the content. These applications depend on the model's ability to capture not just structure, but **semantic understanding**.

In conclusion, semantic modelling is at the heart of modern NLP. It allows machines to go beyond recognizing words and instead grasp the meaning behind them. This progress brings NLP systems closer to human-like language understanding, although many challenges remain. The next sections will explore some of these challenges and how researchers are working to address them.

## 5. Applications and Emerging Use Cases of NLP

As Natural Language Processing (NLP) models have become more powerful and flexible, their use has expanded across many industries. These models are not just used in research—they now support real-world systems that people rely on every day. Below are some of the key areas where NLP is making an impact:

### 1. Healthcare

- **Clinical Summarization:** NLP models help summarize patient notes and medical reports for faster review.
- **Symptom Analysis:** Chatbots and virtual assistants can ask questions and interpret patient symptoms.
- **Biomedical Research:** NLP extracts relevant information from medical journals to support drug discovery and diagnosis.

### 2. Education

- **Essay Evaluation:** NLP can assess grammar, coherence, and structure in student essays.
- **Tutoring Systems:** AI-powered platforms offer language assistance, question answering, and content explanation.
- **Content Summarization:** Large language models can generate simplified summaries of study materials.

### 3. Legal and Finance

- **Contract Analysis:** NLP systems extract clauses, obligations, and conditions from complex legal texts.
- **Sentiment Analysis:** Used in stock market prediction by analysing news, reports, and investor opinions.
- **Document Categorization:** Automatically classifies financial or legal documents by topic or case type.

### 4. Customer Service and Business Intelligence

- **Chatbots:** Handle routine inquiries and support tasks using language models.
- **Feedback Analysis:** NLP systems analyse customer reviews and feedback to identify trends or complaints.
- **Email Triage:** Classifies and routes messages based on urgency or content.

## 5. Media and Communication

- **Content Moderation:** Identifies and flags harmful or offensive content on social platforms.
- **Automated Writing:** Models like GPT generate articles, summaries, and captions.
- **Speech-to-Text Systems:** Used in transcription and voice assistants.

## 6. Emerging Use Cases

- **Multilingual Communication:** Real-time translation and cross-language understanding.
- **Mental Health Support:** NLP tools help analyse speech or writing for emotional cues and early signs of distress.
- **Assistive Technologies:** NLP enables voice-controlled systems for users with disabilities.

### Key Insights:

- NLP is moving from task-specific tools to **general-purpose systems** that adapt across domains.
- Pretrained models reduce the time and cost needed to deploy language technologies.
- As models grow in capability, they are being integrated into **decision-making workflows** - not just language tasks.

## 6. Limitations, Bias, and Ethical Dimensions in NLP

As Natural Language Processing (NLP) systems continue to grow and capability, it is important to understand their limitations and the ethical concerns they raise. While these models achieve impressive results, they are not without flaws. This section outlines the major challenges that need to be addressed to ensure responsible and trustworthy use of NLP.

### 1. Lack of Interpretability

- Many deep learning models, especially large ones like BERT and GPT, are considered “black boxes.”
- It is often unclear how or why a model arrives at a particular output.
- This makes it difficult to trust the model in high-stakes areas like medicine, law, or finance.

### 2. Bias in Language Models

- NLP systems learn from large datasets that may include stereotypes and social biases.
- These biases can appear in model outputs and affect fairness.
- For example, models may link certain professions or behaviours to specific genders or ethnicities.

### 3. Hallucinated and Inaccurate Content

- Generative models can produce convincing but incorrect information.
- This can lead to the spread of misinformation or factual errors.
- In critical domains, such as healthcare, such mistakes can be harmful.

#### **4. Privacy Risks**

- Some models may unintentionally memorize parts of their training data.
- This can lead to the exposure of personal or sensitive information during output generation.
- Ethical use of data and privacy safeguards are essential.

#### **5. Resource and Environmental Concerns**

- Training large language models consumes significant computational power and energy.
- This raises concerns about environmental impact and the accessibility of NLP for smaller institutions or low-resource regions.

Although deep learning has significantly improved NLP capabilities, it also introduces challenges that require ongoing attention. Issues such as bias, interpretability, misinformation, and sustainability must be addressed as part of future research to ensure these systems are ethical, fair, and socially responsible.

#### **7. Future Research Directions in Natural Language Processing**

Natural Language Processing (NLP) has made remarkable progress, but there is still room for improvement. As the field moves forward, researchers are shifting focus toward creating systems that are not only accurate but also fair, interpretable, inclusive, and energy-efficient. Below are some of the key directions that can shape the next generation of NLP technologies.

One important research goal is to improve model interpretability. Most deep learning models today operate as black boxes. Users cannot always understand how a model arrives at a specific decision. This becomes a serious issue in fields like law or medicine, where decisions must be transparent and reliable.

- Future models should offer clearer explanations of how they work.
- Efforts may include visualizations, simplified architectures, or rule-based components added to deep models.

Another area of focus is bias reduction. NLP systems can unintentionally learn and repeat harmful stereotypes from their training data. This can lead to unfair treatment of people based on gender, race, or other attributes.

- Researchers should explore new ways to detect and reduce bias.
- More diverse datasets and fairness-aware algorithms will be essential.

Low-resource and multilingual NLP is also critical. Many languages lack the large datasets that current models require, leaving a gap in accessibility.

- Future models should learn to generalize across languages using smaller datasets.
- Transfer learning and cross-lingual training techniques can help bridge this gap.

Another growing area is multimodal NLP, where models process more than just text. Human communication often involves voice, images, and gestures. Combining these inputs can improve understanding and response quality.

- This can be useful in assistive technologies, education, or customer service.

- Multimodal models will require collaboration across fields like vision, speech, and language.

Finally, sustainability has become a concern. Training large models consumes vast amounts of energy and resources.

- Research should focus on lightweight, efficient architectures.
- Model distillation, pruning, and hardware optimization are promising directions.

In summary, the future of NLP is not just about making models more powerful, but about making them more responsible, inclusive, and accessible. By focusing on these areas, researchers can build language systems that serve a broader range of people and purposes.

## 8. Conclusion

This paper explored the evolution of Natural Language Processing (NLP) from early rule-based systems to today's deep learning-powered language models. By tracing this journey, we highlighted how major shifts in design - such as the move from statistical models to transformer-based architectures—have made it possible for machines to understand, generate, and interact with human language more effectively than ever before.

At the same time, we examined the challenges that come with this progress. Modern NLP systems, while highly capable, are often difficult to interpret, can reflect biases in training data, and require large computational resources. These issues raise important questions about trust, fairness, privacy, and accessibility - especially as NLP becomes embedded in areas like education, law, healthcare, and public communication.

Looking ahead, the future of NLP must be guided by both technical innovation and ethical responsibility. Future research should focus on making models more explainable, reducing social and linguistic bias, supporting low-resource languages, and improving energy efficiency. It is also essential to explore new directions such as multimodal understanding, hybrid systems, and adaptive learning, which can make NLP more interactive and aligned with human needs.

In conclusion, NLP is not only a tool for processing text - it is a powerful technology that shapes how people and machines communicate. Its continued success depends on our ability to balance accuracy with accountability, and progress with purpose. With thoughtful research and collaboration, we can ensure that NLP evolves into a technology that is intelligent, inclusive, and responsible.



## References

1. [Marcos Treviso](#), [Ji-Ung Lee](#), [Tianchu Ji](#), [Betty van Aken](#), [Qingqing Cao](#), [Manuel R. Ciosici](#), [Michael Hassid](#), [Kenneth Heafield](#), [Sara Hooker](#), [Colin Raffel](#), [Pedro H. Martins](#), [André F. T. Martins](#), [Jessica Zosa Forde](#), [Peter Milder](#), [Edwin Simpson](#), [Noam Slonim](#), [Jesse Dodge](#), [Emma Strubell](#), [Niranjan Balasubramanian](#), [Leon Derczynski](#), [Iryna Gurevych](#), [Roy Schwartz](#), “Efficient Methods for Natural Language Processing: A Survey,” 2022. [Online]. Available: <https://arxiv.org/abs/2209.00099>
2. Jonathan J. Webster & Chunyu Kit “Tokenization as the Initial Phase in NLP,” 1992. [Online]. Available: [https://www.researchgate.net/publication/221102283\\_Tokenization\\_as\\_the\\_initial\\_phase\\_in\\_NLP](https://www.researchgate.net/publication/221102283_Tokenization_as_the_initial_phase_in_NLP)
3. Basu Dev Shivahare, Arun Kumar Singh, Nilesh Uppal, Ashar Rizwan, Vangala Sri Vaathsav, Shashikant Suman, “Study of Natural Language Processing and its Recent Applications,” 2022 [Online]. Available: [Survey Paper: Study of Natural Language Processing and its Recent Applications | IEEE Conference Publication | IEEE Xplore](#)
4. Tim Schopf, Karim Arabi and Florian Matthes, “Exploring the Landscape of Natural Language Processing Research,” 2023. [Online]. Available: [https://www.researchgate.net/publication/372487641\\_Exploring\\_the\\_Landscape\\_of\\_Natural\\_Language\\_Processing\\_Research](https://www.researchgate.net/publication/372487641_Exploring_the_Landscape_of_Natural_Language_Processing_Research)
5. Humza Naveed, Asad Ullah Khan, Shi Qiu, Muhammad Saqib, Saeed Anwar, Muhammad Usman, Naveed Akhtar, Nick Barnes, Ajmal Mia, “A Comprehensive Overview of Large Language Models,” 2023. [Online]. Available: [\(PDF\) A Comprehensive Overview of Large Language Models](#)
6. Koroteev M.V., “BERT: A Review of Applications in Natural Language Processing and Understanding,” 2021. [Online]. Available: [https://www.researchgate.net/publication/350287107\\_BERT\\_A\\_Review\\_of\\_Applications\\_in\\_Natural\\_Language\\_Processing\\_and\\_Understanding](https://www.researchgate.net/publication/350287107_BERT_A_Review_of_Applications_in_Natural_Language_Processing_and_Understanding)
7. Venkata Sai Rishita Middi, Middi Appala Raju, Tanvir Ahmed Harris “Machine Translation Using Natural Language Processing,” 2019. [Online]. Available: [\(PDF\) Machine translation using natural language processing](#)
8. Sindhura Kannappan, “Sentiment Analysis Using Natural Language Processing and Machine Learning,” 2023. [Online]. Available: [https://www.researchgate.net/publication/369947834\\_SENTIMENT\\_ANALYSIS\\_USING\\_NATURAL\\_LANGUAGE\\_PROCESSING\\_AND\\_MACHINE\\_LEARNING](https://www.researchgate.net/publication/369947834_SENTIMENT_ANALYSIS_USING_NATURAL_LANGUAGE_PROCESSING_AND_MACHINE_LEARNING)
9. Tom Young , Devamanyu Hazarika , Soujanya Poria , Erik Cambria, “Recent Trends in Deep Learning Based Natural Language Processing,” 2018. [Online]. Available: [1708.02709](#)
10. [Ngoc Tran Khanh Le](#); [Nadia Hadiprodjo](#); [Hazem El-Alfy](#); [Aziz Kerimzhanov](#); [Avtandil Teshebaev](#), “The Recent Large Language Models in NLP,” in 2023 [Online]. Available: [The Recent Large Language Models in NLP | IEEE Conference Publication | IEEE Xplore](#)
11. Saif M. Mohammad, “The State of Natural Language Processing Literature,” 2020. [Online]. Available: [nlpscholar](#)
12. Supriyono, Aji Prasetya Wibawa, Suyono, Fachrul Kurniawan, “Advancements in Natural Language Processing: Implications, Challenges, and Future Directions,” 2024. [Online]. Available: [Advancements in natural language processing: Implications, challenges, and future directions - ScienceDirect](#)
13. [Kai Jiang](#); [Xi Lu](#), “Natural Language Processing and Its Applications in Machine Translation: A Diachronic Review,” in 2020 [Online]. Available: [Natural Language Processing and Its Applications in Machine Translation: A Diachronic Review | IEEE Conference Publication | IEEE Xplore](#)

14. Mohit Mittal, "Understanding Natural Language Processing (NLP) Techniques: From Text Analysis to Language Generation," 2024. [Online]. Available: [\(PDF\) UNDERSTANDING NATURAL LANGUAGE PROCESSING \(NLP\) TECHNIQUES: FROM TEXT ANALYSIS TO LANGUAGE GENERATION](#)
15. Guida Ramos, "Semantic Analyses of Texts for Eliciting and Representing Concepts: The TermCork Project," 2018. [Online]. Available: [\(PDF\) Semantic analyses of texts for eliciting and representing concepts: the TermCork project](#)
16. Akash Arun Kumar Soumya, "Natural Language Processing in Robotics: Leveraging Python for Human-Robot Interaction," 2024. [Online]. Available: [\(PDF\) Natural Language Processing in Robotics: Leveraging Python for Human-Robot Interaction](#)
17. Aditya Jain, Gandhar Kulkarni, Vraj Shah "Natural Language Processing," 2018. [Online]. Available: [https://www.researchgate.net/publication/325772303\\_Natural\\_Language\\_Processing](https://www.researchgate.net/publication/325772303_Natural_Language_Processing)
18. Abu Rayhan, Robert Kinzler, Rajan Rayhan "Natural Language Processing: Transforming How Machines Understand Human Language," 2023. [Online]. Available: [https://www.researchgate.net/publication/373398043\\_NATURAL\\_LANGUAGE\\_PROCESSING\\_TRANSFORMING\\_HOW\\_MACHINES\\_UNDERSTAND\\_HUMAN\\_LANGUAGE](https://www.researchgate.net/publication/373398043_NATURAL_LANGUAGE_PROCESSING_TRANSFORMING_HOW_MACHINES_UNDERSTAND_HUMAN_LANGUAGE)
19. J. Yi; T. Nasukawa; R. Bunescu; W. Niblack, "Sentiment Analyzer: Extracting Sentiments About a Given Topic Using Natural Language Processing Techniques," 2003. [Online]. Available: [Sentiment analyzer: extracting sentiments about a given topic using natural language processing techniques | IEEE Conference Publication | IEEE Xplore](#)
20. Julia Hirschberg and Christopher D. Manning, "Advances in Natural Language Processing," 2017. [Online]. Available: [Advances in natural language processing | Science](#)
21. Tom Young; Devamanyu Hazarika; Soujanya Poria; Erik Cambria, "Recent Trends in Deep Learning Based Natural Language Processing," 2018. [Online]. Available: [Recent Trends in Deep Learning Based Natural Language Processing \[Review Article\] | IEEE Journals & Magazine | IEEE Xplore](#)
22. Scott Miller, Heidi Fox, Lance Ramshaw, Ralph Weischedel, "A Novel Use of Statistical Parsing to Extract Information from Text," 2000[Online]. Available: [A novel use of statistical parsing to extract information from text | Proceedings of the 1st North American chapter of the Association for Computational Linguistics conference](#)
23. Holger Schwenk; Jean-Luc Gauvain, "Connectionist Language Modelling for Large Vocabulary Continuous Speech Recognition," 2002. [Online]. Available: [Connectionist language modeling for large vocabulary continuous speech recognition | IEEE Conference Publication | IEEE Xplore](#)
24. Transductive learning for statistical machine translation, "Proceedings of the 45th Annual Meeting of the Association for Computational Linguistics. [Online]. Available: [Proceedings of the 45th Annual Meeting of the Association of Computational Linguistics](#)